Exercise I

CPSC/AMTH/CBB 663 - Spring semester 2023

Published: Wednesday, January 18, 2023
Due: Friday, February 10, 2023 - <u>11:59 PM</u>

The current problem set requires a working installation of Py-Torch (v1.4), torchvision, numpy, matplotlib, sklearn, and scipy.

Please include all of your written answers and figures in a single PDF, titled `452_<lastname and initals>_assignment1.pdf`, e.g. for a student named Tom Marvolo Riddle, `452_riddletm_assignment1.pdf`. Upload all relevant files for your solutions to Gradescope, including the PDF report and any Python scripts specified. (For this assignment, the Python scripts include `ps1_functions.py` and `prob5_fcnn.py`.) Your homework should be submitted before Friday, February 8, 2023 at 11:59 PM.

We've provided skeleton code for each major function you'll write in a file called `ps1_functions.py`. Please fill in these functions (preserving their names, arguments, and outputs) and include your completed `ps1_functions.py` in your submission (this is needed by our grading scripts). Any supplemental code you write (e.g. calling these functions to generate plots, or trying out different parameters) can be handled however you choose. A well-structured Jupyter notebook with neatly produced and labelled figures is an excellent way to compile assignment reports; just be sure to submit a PDF *and* the separate `ps1_functions.py` file alongside the notebook. However you produce your report, ensure all your figures are clearly labelled.

Programming assignments should use built-in functions in Python and Py-Torch. In general, you may use the `scipy` stack [1]; however, exercises are designed to emphasize the nuances of machine learning and deep learning algorithms - if a function exists that trivially solves an entire problem, please consult with the TA before using it.

Problem 1

Problem 2

Least Squares Solution: $\mathbf{w}^* = (X^\top X)^{-1} X^\top \mathbf{y}$.

$$\textbf{L2 Norm: } \|x\|_2 = \left(\sum_{i=1}^{n} |x_i|^2 \right)^{\frac{1}{2}}$$

1. Write code in Python that randomly generates N points sampled uniformly in the interval $x \in [-1, 3]$. Output the function $y = x^2 - 3x + 1$ for each of the points generated and add zero-mean Gaussian noise with standard deviation σ to y. Make plots of x and y with $N \in \{15, 100\}$ and $\sigma \in \{0, 0.05, 0.2\}$ (there should be six plots in total). Save the point sets for following questions.

   Hint: You may want to check the NumPy library for generating noise.

2. Find the optimal weights (in terms of MSE) for fitting a polynomial function to the data in all 6 cases generated above using a polynomial of degree 1, 2, and 9. Use the least squares analytical solution given above. Do not use built-in methods for regression. Plot the fitted curves on the same plot as the data points (you can plot all 3 polynomial curves on the same plot). Report the fitted weights and the MSE in tables. Qualitatively assess the fit of the curves. Does it look like any of the models overfit, underfit, or appropriately fit the data? Explain your reasoning in one or two sentences.

3. Apply L2 norm regularization with a 9-degree polynomial model to the cases with $\sigma = 0.05$ and $N \in \{15, 100\}$. Vary the parameter λ, and choose three values of λ that result in the following scenarios: underfitting, overfitting, and an appropriate fit. Report the fitted weights and the MSE in each of these scenarios.

   Hint: The least squares solution can also be used for polynomial regression. Check slides of lecture 2 for details on L2 norm regularization.

Problem 3

1. Load the dataset from file `prob3_data_seed.dat` and normalize the features using min-max scaling so that each feature has the same range of values. (The `prob3_data_description.txt` for the description of the provided dataset.)

2. Write a program that applies a k-nn classifier to the data with $k \in \{1, 5, 10, 15\}$. Calculate the test error using both leave-one-out validation and 5-fold cross validation. Plot the test error as a function of k. You may use the existing methods in `scikit-learn` or other libraries for finding the k-nearest neighbors, but do not use any built-in k-nn classifiers. Any reasonable handling of ties in finding

3. Apply two other classifiers of your choice to the same data. For these additional classifiers, you may use existing libraries, such as `scikit-learn` classifiers, but for cross-validation, you should reuse your method from 3.2 or modify it slightly. Possible algorithms include (but are not limited to) logistic regression, QDA, naive Bayes, SVM, and decision trees. Use 5-fold cross validation to calculate the test error. Report the training and test errors. If any tuning parameters need to be selected, use cross-validation and report the training and test error for several values of the tuning parameters. Which of the classifiers performed best? Did any of them underfit or overfit the data? How do they compare to the k-nn classifiers in terms of performance?

Problem 4

1. Suppose we take all the weights and biases in a network of perceptrons, and multiply them by a positive constant, c > 0. Show that the behavior of the network doesn't change. (Exercise in Ch1 Nielsen book)

2. Given the same setup of problem 4.1 - a network of perceptrons - suppose that the overall input to the network of perceptrons has been chosen and fixed. Suppose the weights and biases are such that $w \cdot x + b \neq 0$ for the input x to any particular perceptron in the network. Now replace all the perceptrons in the network by sigmoid neurons, and multiply the weights and biases by a positive constant $c > 0$. Show that in the limit as $c \to \infty$ the behavior of this network of sigmoid neurons is exactly the same as the network of perceptrons. How can this fail when $w \cdot x + b = 0$ for one of the perceptrons? (Exercise in Ch1 Nielsen book)

Use the following figure for problem 4.3 and 4.4.

Figure 1: Multilayer perceptron with three inputs and one hidden layer.

3. For each possible input of the MLP in Figure 1, calculate the output, i.e. what is the output if $X = [0, 0, 0]$, $X = [0, 0, 1]$, etc. You should have 8 cases total.

4. If we change the perceptrons in Figure 1 to sigmoid neurons, what are the outputs for the same inputs (e.g., inputs of $[0, 0, 0], [0, 0, 1], \cdots$)?

5. Using perceptrons with appropriate weights and biases, design an adder that does two-bit binary addition. That is, the adder takes as input two two-bit binary numbers (i.e. 4 binary inputs) and adds them together. Don't forget to include the carry bit. The resulting output should be the two-bit sum and the carry bit for a total of three binary outputs.

Problem 5

1. The time has come — to implement your first fully-connected neural network in PyTorch! For this assignment, we'll be training the network on the canonical MNIST dataset. After building the network, we'll experiment with an array of hyperparameters, tweaking the network's width, depth, learning rate and more in pursuit of the highest classification accuracy we can muster. You may find the Pytorch tutorials helpful as you complete this problem: https://pytorch.org/tutorials/beginner/basics/intro.html. If you haven't yet, we suggest you go through them — especially the tutorial on the optimization loop, which you will need to build more or less from scratch.

We have provided a python template prob5_fcnn.py to implement a neural network with PyTorch based on MNIST data. Within the provided python file is a basic scaffold of a model training workflow. You may use the existing functions in the script for all parts in this problem.

Here are the experiments:

- Follow the TODOs in prob5_fcnn.py to build a two-layer fully-connected neural network. We've provided code to handle the dataset and model initiation, but you need to supply the training logic.
- Try adjusting the learning rate (by making it smaller) if your model is not converging/improving in accuracy. You might also try increasing the number of epochs used.
- Try training your network without a non-linearity between the layers (i.e. a "linear activation"). Then try adding a sigmoid non-linearity, first directly on the input to the first layer, then on the input to the second layer. You should experiment with these independently and in combination. Record your test results for each in a table
- Experiment with the non-linearity used before the middle layer. Here are some activation functions to choose from: relu, softplus, elu, tanh.
- Try changing the width of the hidden layer, keeping the activation function that performs best. Remember to add these results to your table.
- Experiment with the optimizer of your network.
- Lastly, try adding additional layers to your network. How do 3, 4, and 5 layer networks perform? Is there a point where accuracy stops increasing?

Report the results that you got for each experiment. Specifically:

- Provide a single table of the final training and testing errors for all experiments that you have tried.

- Create a plot of the training and testing errors vs the number of iterations for some of the best experiments. How many iterations are sufficient to reach good performance?

2. We will now compare the mean squared error loss and the cross entropy loss. To do this,

    - Finish the provided script to train the model with each of the above loss functions.

    - Create a plot of the training accuracy vs epoch for each loss function (2 lines, 1 plot)

    - Create a plot of the test accuracy vs epoch for each loss function (2 lines, 1 plot)

    Which loss function converges fastest? Which achieves the highest test accuracy? Provide some rational as to the observed differences.

3. Let's now add regularization to the previous network. For the following experiments, you may use the best performing loss function from the previous part. Generalization gap below refers to the (train accuracy - test accuracy).

    - Implement L1 regularization and train the model using $\lambda \in \{0.001, 0.005\}$. Create a plot of the train accuracy, test accuracy, and generalization gap vs epoch for each λ (3 plots, 2 lines each).

    - Implement L2 regularization and train the model using $\lambda \in \{0.001, 0.01, 0.1\}$. Create a plot of the train accuracy, test accuracy, and generalization gap vs epoch for each λ (3 plots, 3 lines each).

    - Apply dropout to both hidden layers and train the model using $p \in \{0.05, 0.1, 0.5\}$. **Hint:** To implement dropout, you can use a special type of PyTorch layer included in `torch.nn` [2]. Create a plot of the train accuracy, test accuracy, and generalization gap vs epoch for each p value (3 plots, 3 lines each)

    - Using the loss data you've collected so far, create a plot of the test accuracy vs epoch for each of the experiments performed for this section. (8 lines, 1 plot)

    Are the final results sensitive to each parameter? Is there any regularization method which performs best?

4. Print a confusion matrix showing which digits were misclassified, and what they were misclassified as. What numbers are frequently confused with one another by your model? (You may use `sklearn`'s confusion matrix function to generate the matrix.)

5. What was the highest percentage of classification accuracy your fully-connected network achieved? Briefly describe the architecture and training process that produced it. (If you like, you can take part in our friendly class competition by posting your results, along with a short description of your methods, to https://edstem.org/us/courses/33793/discussion/2396501.)

References

[1] "The scipy stack specification." [Online]. Available: https://www.scipy.org/stackspec.html

[2] "Pytorch nn module docs." [Online]. Available: https://pytorch.org/docs/stable/nn.html