

20230410_02_word2vec

April 16, 2023

0.1 # Problem 2 Word2Vec & Seq2Seq Models

0.2 Setup

```
[ ]: import os, sys, tqdm, json, pickle, random, warnings
      from typing import List, Optional

      import numpy as np, pandas as pd

      import matplotlib as mpl, matplotlib.pyplot as plt, seaborn as sns

      import torch.nn.functional as F

      import torch, torch.nn as nn, torch.optim as optim
      import pytorch_lightning as pl
      from torch.utils import data

      from gensim.models import Word2Vec, Phrases
      from gensim.test.utils import common_texts
      import gensim, gensim.downloader as gensim_api
      from gensim import corpora

      from gensim.parsing.preprocessing import remove_stopwords, preprocess_documents
      from gensim.utils import simple_preprocess

      pl.seed_everything(3)
```

Global seed set to 3

```
[ ]: 3
```

0.2.1 Global Variables

NOTE: you can find other gensim datasets [here](#)

```
[ ]: SAVE_DIR = os.path.abspath('.')

      W2V_FILE = os.path.join(SAVE_DIR, 'word2vec.pth')
```

```

CORPUS = 'text8'# 'glove-wiki-gigaword-200'

TXT_FILE = os.path.join(SAVE_DIR, f'{CORPUS}.pkl')

OVERWRITE = False

VEC_SIZE = 50 #10

#DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
DEVICE = 'cpu'

SAVE_DIR

```

```
[ ]: '/Users/default/Google Drive/currentDocumants/studies/Master/6.Semester/deep_learning/pset/pset3/problem 2'
```

0.3 Part 1

0.3.1 Fetch you text corpus

TODO: generate your own word-vector embedding using a text corpus of *your* choice of at least **500** words. You may use `gensim`'s corpora.

Feel free to play with hyperparamters in the notebook such as context size and embedding size.

```

[ ]: corpus = gensim_api.load(CORPUS)

[ ]: does_corpus_exist = os.path.isfile(TXT_FILE)

if not does_corpus_exist or OVERWRITE:
    corpus = gensim_api.load(CORPUS)
    with open(TXT_FILE, 'wb') as f:
        pickle.dump(corpus, f)

else:
    with open(TXT_FILE, 'rb') as f:
        corpus = pickle.load(f)

corpus

```

```
[ ]: <text8.Dataset at 0x7f8fdaafa200>
```

0.3.2 Train your Word2Vec Model

NOTE: you are free to create your own Word2Vec model. You are welcome to use `gensim`. However for the pset you must at least *train* a Word2Vec model!

```
[ ]: W2V_FILE
```

```
[ ]: '/Users/default/Google Drive/currentDocumants/studies/Master/6.Semester/deep_learning/pset/pset3/problem 2/word2vec.pth'
```

```
[ ]: does_model_exist = os.path.isfile(W2V_FILE)

if not does_model_exist or OVERWRITE:
    w2v = Word2Vec(
        # NOTE: you can use either a list of strings (sentences)
        # or provide a corpus_file
        corpus,

        # Dimensionality of the word vectors.
        vector_size=VEC_SIZE,
        #size=VEC_SIZE,

        # Maximum distance between the current and predicted word within a
        ↪sentence.
        window=5,

        # Ignores all words with total frequency lower than this.
        min_count=1,

        # Training algorithm: 1 for skip-gram; otherwise CBOW.
        sg = 1,

        # Use these many worker threads to train the model (=faster training
        # with multicore machines).
        workers=32,

        # If 0, use the sum of the context word vectors. If 1, use the mean,
        # only applies when cbow is used.
        cbow_mean = 1,

        # Number of iterations (epochs) over the corpus. (Formerly: iter)
        epochs=10#100
    )

    w2v.add_null_word()

    w2v.save(W2V_FILE)

else:

    w2v = Word2Vec.load(W2V_FILE)
```

```
[ ]: def vec_math(model, *args, return_str_description:bool=False):
    desc = ''
```

```

func = np.add
# NOTE: vector_size
vect = None

for arg in args:
    if arg == '+':
        desc += ' + '
        func = np.add
    elif arg == '-':
        desc += ' - '
        func = np.subtract
    elif arg == '/':
        desc += ' / '
        func = np.divide
    elif arg == '*':
        desc += ' * '
        func = np.multiply
    else:
        desc += arg
        curr = model.wv[arg]
        if vect is None:
            vect = curr.copy()
        else:
            vect = func(vect, curr)

if return_str_description:
    return desc, vect

return vect

```

Explore results NOTE: here my model had

$$woman + man - king = girl$$

and

$$news + truth - lies = blogging$$

```

[ ]: # woman + man - king = ?
gsm_res = w2v.wv.most_similar(positive=['woman', 'man'], negative=['king'])[0]
gsm_res

```

```

[ ]: ('dumb', 0.7257739901542664)

```

```
[ ]: # news + truth - lies = ?
gsm_res = w2v.wv.most_similar(positive=['news', 'truth'], negative=['lies'])[0]
gsm_res
```

```
[ ]: ('publicize', 0.7525402903556824)
```

```
[ ]: w2v.wv.most_similar(positive=['lies'], topn = 10)
```

```
[ ]: [('extends', 0.8267737030982971),
      ('occupies', 0.8245770335197449),
      ('lying', 0.8212238550186157),
      ('situated', 0.8177280426025391),
      ('reaches', 0.8083112835884094),
      ('encircles', 0.766158938407898),
      ('traverses', 0.7654803395271301),
      ('sits', 0.762751579284668),
      ('overlooks', 0.7599471211433411),
      ('encloses', 0.7521087527275085)]
```

Visualize **TODO:** Plot these word embeddings (all words) by adjusting the code in the notebook.

You are able to adjust the plotting parameters to suit your needs for making a compelling visualization. Discuss what you notice in your embeddings. For example, using the introduction to Charles Darwin’s “*On the Origin of Species*” as a text file, we obtain the embeddings in Figure 2

```
[ ]: corpus
```

```
[ ]: <text8.Dataset at 0x7f8fdaafa200>
```

```
[ ]: #extract words from corpus
all_words_from_corpus = []
for word in corpus:
    all_words_from_corpus.append(word)

#words_from_corpus = np.array(words_from_corpus).flatten()
all_words_from_corpus = np.concatenate(all_words_from_corpus)
```

```
[ ]: #this does not terminate

#fig = plt.figure(figsize=(20, 8))
#lim_scale = 2

#ax = fig.add_subplot(1, 1, 1)

# Get the 50 most frequent words in the vocabulary
#words = [w for w, _ in w2v.wv.most_similar(positive=['king'], topn=20)]
```

```

#vecs = w2v.wv[all_words_from_corpus]

#min_vals = np.min(vecs, axis=0)
#max_vals = np.max(vecs, axis=0)

#ax.set_xlim(min_vals[0] * lim_scale, max_vals[0] * lim_scale)
#ax.set_ylim(min_vals[1] * lim_scale, max_vals[1] * lim_scale)

#for j, vec in enumerate(vecs):
#    word = all_words_from_corpus[j]

#    #print(vec)
#    plt.scatter(vec[0], vec[1])
#    plt.annotate(
#        word, xy=(0, 0), xytext=(vec[0], vec[1]),
#        arrowprops=dict(arrowstyle="<-"),
#        fontsize=8
#    )

#plt.show()

```

TODO: Since this is too crowded to interpret, modify the code in the notebook to randomly select words to plot as long as there is space as shown in Figure 3

```
[ ]: words_sample = np.random.choice(all_words_from_corpus, size=20, replace=False)
```

```
[ ]: fig = plt.figure(figsize=(20, 8))
lim_scale = 2

ax = fig.add_subplot(1, 1, 1)

vecs = w2v.wv[words_sample]

min_vals = np.min(vecs, axis=0)
max_vals = np.max(vecs, axis=0)

ax.set_xlim(min_vals[0] * lim_scale, max_vals[0] * lim_scale)
ax.set_ylim(min_vals[1] * lim_scale, max_vals[1] * lim_scale)

for j, vec in enumerate(vecs):
    word = words_sample[j]

    #print(vec)
    plt.scatter(vec[0], vec[1])
    plt.annotate(
        word, xy=(0, 0), xytext=(vec[0], vec[1]),
        arrowprops=dict(arrowstyle="<-"),


```

```

        fontsize=8 # Adjust the font size for readability
    )

plt.show()

```



subset, todo delete

```

[ ]: fig = plt.figure(figsize=(20, 8))
lim_scale = 2

word_groups = [
    'computer,news,truth,lies,big,news + truth - lies'.split(','),
    'man,king,woman,queen,woman + man - king'.split(','),
]

for i, words in enumerate(word_groups):
    ax = fig.add_subplot(1, 2, i+1)
    single_words = list(filter(lambda e: e.isalpha(), words))
    domath_words = list(filter(lambda e: not e.isalpha(), words))

    vecs = w2v.wv[single_words]
    for w in domath_words:
        v = vec_math(w2v, *w.split())
        vecs = np.vstack((vecs, v))

    min_vals = np.min(vecs, axis=0)
    max_vals = np.max(vecs, axis=0)

    ax.set_xlim(min_vals[0] * lim_scale, max_vals[0] * lim_scale)
    ax.set_ylim(min_vals[1] * lim_scale, max_vals[1] * lim_scale)


```

```

for j, vec in enumerate(vecs):
    word = words[j]

    plt.scatter(vec[0], vec[1])
    plt.annotate(
        word, xy=(0, 0), xytext=(vec[0], vec[1]),
        arrowprops=dict(arrowstyle="<-")
    )

```



0.4 Part 2

0.4.1 Dataset

dl_310_2

```

[ ]: import os, sys, tqdm, json, pickle, random, warnings
      from typing import List, Optional

      import numpy as np, pandas as pd

      import matplotlib as mpl, matplotlib.pyplot as plt, seaborn as sns

      import torch, torch.nn as nn, torch.optim as optim
      import pytorch_lightning as pl
      from torch.utils import data

      from gensim.models import Word2Vec, Phrases
      from gensim.test.utils import common_texts
      import gensim, gensim.downloader as gensim_api
      from gensim import corpora

      from gensim.parsing.preprocessing import remove_stopwords, preprocess_documents

```



```
from gensim.utils import simple_preprocess

pl.seed_everything(3)
```

Global seed set to 3

[]: 3

Whether or not you used `gensim` for training the Word2Vec model you will be using PyTorch Lightning for the Seq2Seq model. To make this easier on you, you will need to create a `data.Dataset` object which contains, at the very least, `__getitem__` and `__len__` methods to yield items during training. Below is an example which converts a `gensim` corpus paired batches of the first 300 words per doc.

```
[ ]: class Seq2SeqTextDataset(data.Dataset):
    def __init__(self, corpus, keyed_vecs, reverse_target:bool=False):
        self.corpus = corpus
        self.keyed_vecs = keyed_vecs

        docs = self.docs_from_corpus(corpus)
        self.docs = docs
        self.data = docs

        corpora_dct = corpora.Dictionary(docs)
        self.corpora_dct = corpora_dct

        bows = self.bows_from_docs(docs, corpora_dct)
        self.bows = bows

        self.reverse_target = reverse_target

    def docs_from_corpus(self, corpus):
        docs = [
            #simple_preprocess(remove_stopwords(' '.join(doc)))[:300] #will
            → have to reduce this to 5
            simple_preprocess(remove_stopwords(' '.join(doc)))[:3]
            for doc in corpus
        ]
        return docs

    def bows_from_docs(self, docs, dct):
        return [dct.doc2bow(doc) for doc in docs]

    def __len__(self):
        return int(len(self.docs))
```

```

def __getitem__(self, idx):
    doc = self.docs[idx]

    source = torch.tensor([self.keyed_vecs.get_index(word) for word in doc])

    target = [self.keyed_vecs.get_index(word, default=0) for word in doc]
    if self.reverse_target:
        target = target[::-1]
    target = torch.tensor(target)

    return source, target

```

0.4.2 Seq2Seq Model

Encoder TODO: fill in the encode method

```

[ ]: class Encoder(nn.Module):
    def __init__(
        self,
        n_vocabsize:int, n_embedding:int, n_hidden:int, n_layers:int,
        dropout:Optional[float]=0.1, pretrained:Optional[np.ndarray]=None,
    ):
        super(Encoder, self).__init__()

        self.n_vocabsize = n_vocabsize
        self.n_embedding = n_embedding
        self.n_hidden = n_hidden
        self.n_layers = n_layers
        self.dropout = dropout

        embedding = self.make_embedding_layer(n_vocabsize, n_embedding,
↳pretrained)
        self.embedding = embedding

        # NOTE: depending on your dataset you may have to change batch_first
        self.lstm = nn.LSTM(n_embedding, n_hidden, n_layers, dropout=dropout,
↳batch_first=True)
        self.dropout = nn.Dropout(dropout)

    def make_embedding_layer(
        self,
        n_vocabsize:int,
        n_embedding:int,
        pretrained:Optional[np.ndarray]=None,
    ):
        if pretrained is not None:
            pretrained = np.array(pretrained)

```

```

        n_vocabsize, n_embedding = pretrained.shape

        self.n_vocabsize = n_vocabsize
        self.n_embedding = n_embedding

        pretrained = torch.FloatTensor(pretrained)
        embedding = nn.Embedding.from_pretrained(pretrained)

    else:
        embedding = nn.Embedding(n_vocabsize, n_embedding)

    return embedding

def encode(self, x):
    # TODO: fill this in
    embedding = self.embedding(x)
    embedding = self.dropout(embedding) #TODO: is this the right place for
    ↪ dropout?

    _, (hidden, cell) = self.lstm(embedding) #TODO save output?
    return hidden, cell

def forward(self, x):
    return self.encode(x)

```

Decoder **TODO:** fill in the decode method

```

[ ]: class Decoder(nn.Module):
    def __init__(
        self,
        n_vocabsize:int, n_embedding:int, n_hidden:int, n_layers:int,
        dropout:Optional[float]=0.1,
    ):
        super(Decoder, self).__init__()
        self.n_vocabsize = n_vocabsize
        self.n_embedding = n_embedding
        self.n_hidden = n_hidden
        self.n_layers = n_layers
        self.dropout = dropout

        self.embedding = nn.Embedding(n_vocabsize, n_embedding)
        self.lstm = nn.LSTM(n_embedding, n_hidden, n_layers, dropout=dropout,
    ↪ batch_first=True) #TODO batch_first?

        self.fc_out = nn.Linear(n_hidden, n_vocabsize)
        self.dropout = nn.Dropout(dropout)

```

```

def decode(self, x, hidden, cell):
    x = x.unsqueeze(1)
    embedded = self.embedding(x)

    output, (hidden, cell) = self.lstm(embedded, (hidden, cell))

    prediction = self.fc_out(output.squeeze(1))
    prediction = self.dropout(prediction) #TODO: is this the right place_
    ↪ for dropout?

    return prediction, hidden, cell

def forward(self, x, hidden, cell):
    return self.decode(x, hidden, cell)

```

Seq2Seq LightningModule **TODO:** fill in the do_seq2seq method and define the criterion

```

[ ]: class Seq2Seq(pl.LightningModule):
    def __init__(
        self,
        word_2_vec,
        train_loader,
        n_vocabsize:int=0,
        n_hidden:int = 2,
        n_layers:int = 2,
        dropout:float = 0.2,
        teacher_forcing_ratio:float=0.5,
        learning_rate:float=0.01,
    ):

        super(Seq2Seq, self).__init__()

        self.w2v = word_2_vec
        self.train_loader = train_loader

        n_embedding = word_2_vec.vector_size
        n_vocabsize, n_embedding = np.array(word_2_vec.wv).shape

        self.n_vocabsize = n_vocabsize
        self.n_embedding = n_embedding
        self.n_hidden = n_hidden
        self.n_layers = n_layers
        self.n_output = n_vocabsize

```

```

self.teacher_forcing_ratio = teacher_forcing_ratio
self.learning_rate = learning_rate

self.encoder = Encoder(
    n_vocabsize, n_embedding, n_hidden, n_layers,
    dropout, pretrained=np.array(word_2_vec.wv)
)
self.decoder = Decoder(
    n_vocabsize, n_embedding, n_hidden, n_layers, dropout,
)

self.criterion = torch.nn.CrossEntropyLoss() #TODO: is this the right
↪ loss function?

def train_dataloader(self):
    return self.train_loader

def configure_optimizers(self):
    optimizer = torch.optim.AdamW(self.parameters(), lr=self.learning_rate,
↪ weight_decay=1e-5)
    scheduler = optim.lr_scheduler.StepLR(optimizer, 100, gamma=0.99)
    return [optimizer], [scheduler]

def encode(self, word):
    index = self.w2v.wv.get_index(word)
    return torch.tensor(index)

def do_seq2seq(self, source, target):
    # TODO: fill this in
    batch_size = source.shape[0]
    target_len = target.shape[1] #TODO 1 or 0?
    target_vocab_size = self.n_vocabsize

    outputs = torch.zeros(target_len, batch_size, target_vocab_size).
↪ to(self.device)

    hidden, cell = self.encoder(source)

    x = target[0]

    for t in range(1, target_len):
        output, hidden, cell = self.decoder(x, hidden, cell)

        outputs[t] = output

        teacher_force = random.random() < self.teacher_forcing_ratio

```

```

        top1 = output.argmax(1)

        x = target[t] if teacher_force else top1

    return outputs

def training_step(self, batch, batch_idx):
    x, y = batch
    source = x
    target = y

    output = self.do_seq2seq(source, target)

    # NOTE: this may not be needed depending on your
    output_dim = output.shape[-1]
    output = output.view(-1, output_dim)
    target = target.view(-1)

    loss = self.criterion(output, target)

    result = {'loss': loss}
    self.log('loss', loss)
    return result

```

```

[ ]: import torch.nn.functional as F

class Seq2Seq(pl.LightningModule):
    def __init__(
        self,
        word_2_vec,
        train_loader,
        n_vocabsize:int=0,
        n_hidden:int = 2,
        n_layers:int = 2,
        dropout:float = 0.2,
        teacher_forcing_ratio:float=0.5,
        learning_rate:float=0.01,
    ):

        super(Seq2Seq, self).__init__()

        self.w2v = word_2_vec
        self.train_loader = train_loader

        n_embedding = word_2_vec.vector_size

```

```

n_vocabsize, n_embedding = word_2_vec.wv.vectors.shape
#n_vocabsize, n_embedding = word_2_vec.shape

self.n_vocabsize = n_vocabsize
self.n_embedding = n_embedding
self.n_hidden = n_hidden
self.n_layers = n_layers
self.n_output = n_vocabsize

self.teacher_forcing_ratio = teacher_forcing_ratio
self.learning_rate = learning_rate

self.encoder = Encoder(
    n_vocabsize, n_embedding, n_hidden, n_layers,
    dropout, pretrained=word_2_vec.wv.vectors # <-- Pass 'vectors'
↳attribute
)

#self.encoder = Encoder(
#    n_vocabsize, n_embedding, n_hidden, n_layers,
#    dropout, pretrained=np.array(word_2_vec.wv))

self.decoder = Decoder(
    n_vocabsize, n_embedding, n_hidden, n_layers, dropout,
)

self.criterion = torch.nn.CrossEntropyLoss()

def train_dataloader(self):
    return self.train_loader

def configure_optimizers(self):
    optimizer = torch.optim.AdamW(self.parameters(), lr=self.learning_rate,
↳weight_decay=1e-5)
    scheduler = optim.lr_scheduler.StepLR(optimizer, 100, gamma=0.99)
    return [optimizer], [scheduler]

def encode(self, word):
    index = self.w2v.wv.get_index(word)
    return torch.tensor(index)

def do_seq2seq(self, source, target):
    batch_size = source.shape[0]
    trg_len = target.shape[1]
    trg_vocab_size = self.n_output

```

```

        outputs = torch.zeros(trg_len, batch_size, trg_vocab_size).to(self.
↪device)

        hidden, cell = self.encoder(source)

        input = target[:, 0]

        for t in range(1, trg_len):
            #output, hidden = self.decoder(input, hidden, encoder_outputs)
            #output, (hidden, cell) = self.decoder(input, hidden, cell)
            output, hidden, cell = self.decoder(input, hidden, cell)
            outputs[t] = output
            teacher_force = random.random() < self.teacher_forcing_ratio
            top1 = output.argmax(1)
            input = target[:, t] if teacher_force else top1

        return outputs

def training_step(self, batch, batch_idx):
    x, y = batch
    source = x
    target = y

    output = self.do_seq2seq(source, target)

    output_dim = output.shape[-1]
    output = output.view(-1, output_dim)
    target = target.view(-1)

    loss = self.criterion(output, target)

    result = {'loss': loss}
    self.log('loss', loss)
    return result

```

0.4.3 Instantiate dataset and dataloader

```
[ ]: ds = Seq2SeqTextDataset(corpus, w2v.wv)
```

```
[ ]: dl = data.DataLoader(ds, batch_size=3, shuffle=True, drop_last=True)
```


0.4.4 Instantiate model

```
[ ]: s2s = Seq2Seq(
    w2v, dl,
    n_hidden=12, n_layers=4, dropout=0.5, learning_rate = 0.01
).to(DEVICE)
s2s.encoder.to(DEVICE)
s2s.decoder.to(DEVICE)
s2s
```

```
[ ]: Seq2Seq(
  (encoder): Encoder(
    (embedding): Embedding(253854, 50)
    (lstm): LSTM(50, 12, num_layers=4, batch_first=True, dropout=0.5)
    (dropout): Dropout(p=0.5, inplace=False)
  )
  (decoder): Decoder(
    (embedding): Embedding(253854, 50)
    (lstm): LSTM(50, 12, num_layers=4, batch_first=True, dropout=0.5)
    (fc_out): Linear(in_features=12, out_features=253854, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
  )
  (criterion): CrossEntropyLoss()
)
```

0.4.5 Define trainer

```
[ ]: trainer = pl.Trainer(
    max_epochs=4,

    # NOTE: gradient clipping can help prevent exploding gradients
    gradient_clip_val=100,
    gradient_clip_algorithm='value',
    log_every_n_steps=5,

    # NOTE: this should match your device i.e. if you set cuda above, this
    ↪ should be cuda.
    # Otherwise it should be cpu.
    accelerator=DEVICE,

    # NOTE: you can set the maximum time you want to train your model
    max_time={'minutes': 5},

    # NOTE: setting this to true will save your model every so often
    enable_checkpointing=False,
    accumulate_grad_batches=2
)
```

```
GPU available: False, used: False
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs
```

0.4.6 Train model

TODO train your model

```
[ ]: trainer.fit(s2s, dl)
```

```

| Name      | Type           | Params
-----
0 | encoder   | Encoder        | 12.7 M
1 | decoder   | Decoder        | 16.0 M
2 | criterion | CrossEntropyLoss | 0
-----
16.0 M    Trainable params
12.7 M    Non-trainable params
28.7 M    Total params
114.797   Total estimated model params size (MB)
/Users/default/miniconda3/envs/dl_310_2/lib/python3.10/site-
packages/pytorch_lightning/trainer/connectors/data_connector.py:430:
PossibleUserWarning: The dataloader, train_dataloader, does not have many
workers which may be a bottleneck. Consider increasing the value of the
`num_workers` argument` (try 8 which is the number of cpus on this machine) in
the `DataLoader` init to improve performance.
  rank_zero_warn(
Epoch 0: 100%|          | 567/567 [02:03<00:00,  4.59it/s, v_num=11]
`Trainer.fit` stopped: `max_epochs=1` reached.
Epoch 0: 100%|          | 567/567 [02:03<00:00,  4.59it/s, v_num=11]
```

TODO: Train your network also to reconstruct sentences with some words masked. How well does this training perform? Report accuracy at convergence.

0.4.7 Evaluation

```
[ ]: for source, target in dl:
      break
```

```
[ ]: res = s2s.do_seq2seq(
      source[0, :].unsqueeze(0),
      target[0, :].unsqueeze(0)
```

```
)
res = res.squeeze(1)
res.shape
```

```
[ ]: torch.Size([3, 253854])
```

```
[ ]: [w2v.wv.index_to_key[idx] for idx in res.argmax(1)]
```

```
[ ]: ['the', 'held', 'seven']
```

TODO: Put in a sentence from the corpus and write its reconstruction. Is the reconstruction perfect?

```
[ ]: #example_sentence = 'the seven kings'
```

```
[ ]: example_sentence = all_words_from_corpus[1:20]
```

```
[ ]: tokens = example_sentence
tokens = [w2v.wv.get_index(token, default=0) for token in tokens]

tokens = torch.tensor(tokens).unsqueeze(0)
tokens
```

```
[ ]: tensor([[ 3082,    11,     5,   194,     1,  3133,    45,    58,   155,   127,
            741,   476, 10578,   133,     0, 27455,     1,     0,   102]])
```

```
[ ]: # Set the target tensor to be the same as the tensor
target_tensor = tokens.clone()
target_tensor
```

```
[ ]: tensor([[ 3082,    11,     5,   194,     1,  3133,    45,    58,   155,   127,
            741,   476, 10578,   133,     0, 27455,     1,     0,   102]])
```

```
[ ]: random.seed(1)
# Generate the output tensor using the seq2seq model
output_tensor = s2s.do_seq2seq(tokens, target_tensor).argmax(2)
output_tensor = output_tensor.squeeze(1).numpy()
output_tensor
```

```
[ ]: array([[    0,   5626,  76185,   150,   3782,    22,   126,    22,
            76185,   486,   938,   359,    22,  3429,    22,   486,
            247, 225309,    22])
```

```
[ ]: # Convert the output tensor back to words
#ds.keyed_vecs.index_to_key[int(22)]
reconstructed_sentence = [ds.keyed_vecs.index_to_key[int(idx)] for idx in
    ↪output_tensor]
```

```
reconstructed_sentence
```

```
[ ]: ['the',  
      'peaceful',  
      'miklos',  
      'french',  
      'obsolete',  
      'seven',  
      'university',  
      'seven',  
      'miklos',  
      'middle',  
      'unit',  
      'held',  
      'seven',  
      'watch',  
      'seven',  
      'middle',  
      'word',  
      'acetabularia',  
      'seven']
```

```
[ ]: print("Original sentence:")  
     print(" ".join(example_sentence))  
  
     print("\nReconstructed sentence:")  
     print(" ".join(reconstructed_sentence))
```

Original sentence:

originated as a term of abuse first used against early working class radicals including the diggers of the english

Reconstructed sentence:

the peaceful miklos french obsolete seven university seven miklos middle unit held seven watch seven middle word acetabularia seven

TODO: Now mask one of the words in the sentence with xxxx and test if it fills the word back.

```
[ ]: from random import randint  
  
masked_sentence = example_sentence.copy()  
#get the index of the word to be masked  
masked_index = random.randint(0, len(masked_sentence)-1)  
#replace the word with the mask token  
masked_sentence[mask_index] = 'xxxx'  
masked_sentence
```

```
[ ]: array(['originated', 'as', 'a', 'term', 'of', 'abuse', 'first', 'used',
          'against', 'early', 'xxxx', 'class', 'radicals', 'including',
          'the', 'diggers', 'of', 'the', 'english'], dtype='<U100')
```

```
[ ]: # Convert the masked sentence to a tensor
masked_tokens = [w2v.wv.get_index(token, default=0) for token in
↳masked_sentence]
masked_tokens = torch.tensor(masked_tokens).unsqueeze(0)
masked_tokens
```

```
[ ]: tensor([[ 3082,    11,     5,   194,     1,  3133,    45,    58,   155,   127,
              30522,   476, 10578,   133,     0, 27455,     1,     0,   102]])
```

```
[ ]: # Set the target tensor to be the same as the tensor
target_tensor = masked_tokens.clone()
target_tensor
```

```
[ ]: tensor([[ 3082,    11,     5,   194,     1,  3133,    45,    58,   155,   127,
              30522,   476, 10578,   133,     0, 27455,     1,     0,   102]])
```

```
[ ]: #reconstruct the masked sentence
output_tensor = s2s.do_seq2seq(masked_tokens, target_tensor).argmax(2)
output_tensor = output_tensor.squeeze(1).numpy()
output_tensor
```

```
[ ]: array([[ 0,    85,   801, 2347,   247,    7,   298,   22,   22,
            4334,   22,   486,    7,    7,  4334,   22,   22, 12578,
            4334])
```

```
[ ]: # Convert the output tensor back to words
reconstructed_sentence = [ds.keyed_vecs.index_to_key[int(idx)] for idx in
↳output_tensor]
reconstructed_sentence
```

```
[ ]: ['the',
      'united',
      'leading',
      'fleet',
      'word',
      'zero',
      'best',
      'seven',
      'seven',
      'explosion',
      'seven',
      'middle',
      'zero',
```

```
'zero',
'explosion',
'seven',
'seven',
'slit',
'explosion']
```

```
[ ]: print("Original sentence:")
      print(" ".join(example_sentence))

      print("Masked sentence:")
      print(" ".join(masked_sentence))

      print("\nReconstructed sentence:")
      print(" ".join(reconstructed_sentence))
```

Original sentence:

originated as a term of abuse first used against early working class radicals including the diggers of the english

Masked sentence:

originated as a term of abuse first used against early xxxx class radicals including the diggers of the english

Reconstructed sentence:

the united leading fleet word zero best seven seven explosion seven middle zero zero explosion seven seven slit explosion

TODO: Finally, start a sentence from the corpus by giving the first 3 words, with the rest of the words masked, see if it completes this sentence.

```
[ ]: masked_sentence = example_sentence.copy()

      #replace all words but the first three with the mask token
      for i in range(3, len(masked_sentence)):
          masked_sentence[i] = 'xxxx'

      masked_sentence
```

```
[ ]: array(['originated', 'as', 'a', 'xxxx', 'xxxx', 'xxxx', 'xxxx', 'xxxx',
          'xxxx', 'xxxx', 'xxxx', 'xxxx', 'xxxx', 'xxxx', 'xxxx', 'xxxx',
          'xxxx', 'xxxx', 'xxxx'], dtype='<U100')
```

```
[ ]: # Convert the masked sentence to a tensor
      masked_tokens = [w2v.wv.get_index(token, default=0) for token in
          ↪masked_sentence]
      masked_tokens = torch.tensor(masked_tokens).unsqueeze(0)
      masked_tokens
```

```
[ ]: tensor([[ 3082,    11,     5, 30522, 30522, 30522, 30522, 30522, 30522, 30522,
              30522, 30522, 30522, 30522, 30522, 30522, 30522, 30522]])
```

```
[ ]: # Set the target tensor to be the same as the tensor
target_tensor = masked_tokens.clone()
target_tensor
```

```
[ ]: tensor([[ 3082,    11,     5, 30522, 30522, 30522, 30522, 30522, 30522, 30522,
              30522, 30522, 30522, 30522, 30522, 30522, 30522, 30522]])
```

```
[ ]: #reconstruct the masked sentence
output_tensor = s2s.do_seq2seq(masked_tokens, target_tensor).argmax(2)
output_tensor = output_tensor.squeeze(1).numpy()
output_tensor
```

```
[ ]: array([[ 0, 801, 76185, 486, 1342, 22, 359, 22, 22,
             3782, 938, 22, 22, 22, 4334, 22, 486, 176,
             22])
```

```
[ ]: # Convert the output tensor back to words
reconstructed_sentence = [ds.keyed_vecs.index_to_key[int(idx)] for idx in
↪output_tensor]
reconstructed_sentence
```

```
[ ]: ['the',
      'leading',
      'miklos',
      'middle',
      'format',
      'seven',
      'held',
      'seven',
      'seven',
      'seven',
      'obsolete',
      'unit',
      'seven',
      'seven',
      'seven',
      'seven',
      'explosion',
      'seven',
      'middle',
      'common',
      'seven']
```

```
[ ]: print("Original sentence:")
print(" ".join(example_sentence))
```

```
print("Masked sentence:")
print(" ".join(masked_sentence))

print("\nReconstructed sentence:")
print(" ".join(reconstructed_sentence))
```

Original sentence:

originated as a term of abuse first used against early working class radicals
including the diggers of the english

Masked sentence:

originated as a xxxx xxxx xxxx xxxx xxxx xxxx xxxx xxxx xxxx xxxx xxxx
xxxx xxxx xxxx

Reconstructed sentence:

the leading miklos middle format seven held seven seven obsolete unit seven
seven seven explosion seven middle common seven

0.5 Exploration

You can also use `gensim` to load other models aside from `Word2Vec` such as `glove`

```
[ ]: #glove_vectors = gensim_api.load('glove-twitter-25')
      #glove_vectors.most_similar('twitter')
```

0.6 Citations

This notebook adapts parts of [@Ben Trevett's PyTorch Seq2Seq notebook](#). Additionally we utilize parts of both PyTorch's [documentation](#) as well as [Gensim's](#) and [Machine Learning Plus's tutorial](#)

0.7 Supplementary materials

For those getting started with [gensim](#) they may find [Gaurav Padawe tutorial](#) useful. Of course there is also the official `Word2Vec` [gensim tutorial](#)

For support with [Pylightning](#) please refer to these documentation pages: - [Trainer documentation](#)
- [LightningModule documentation](#)